

## Class\_08

Anika Bhattacharjya (A15459876)

10/21/2021

First up is clustering methods

#Kmeans clustering

The function in base R to do Kmeans clustering is called 'kmeans()'


First make up some data where we know what the answer should be:

```
tmp <- c(rnorm(30,-3), rnorm(30,3))
x <- cbind(x=tmp, y=rev(tmp))
x
```

```
##           x           y
## [1,] -3.903877 0.359280
## [2,] -2.521531 2.976836
## [3,] -3.284051 4.020068
## [4,] -3.883268 3.298901
## [5,] -2.936833 2.123831
## [6,] -4.130707 3.751103
## [7,] -1.836739 5.196439
## [8,] -2.475593 1.375293
## [9,] -4.114520 2.961445
## [10,] -5.056103 2.106063
## [11,] -3.964889 2.529022
## [12,] -2.631064 2.817329
## [13,] -2.776078 4.867017
## [14,] -3.211672 3.049036
## [15,] -1.918888 4.484413
## [16,] -3.584845 4.432769
## [17,] -3.232274 1.271517
## [18,] -2.312825 3.448025
## [19,] -4.269171 2.349272
## [20,] -2.475074 3.866457
## [21,] -2.658771 2.088574
## [22,] -1.568452 1.159585
## [23,] -3.656608 2.987012
## [24,] -4.374141 2.677371
## [25,] -2.522436 1.439703
## [26,] -2.673494 2.333662
## [27,] -2.518130 2.645605
## [28,] -4.468089 1.289056
## [29,] -2.165905 4.639808
## [30,] -3.300228 5.059465
```

```
## [31,] 5.059465 -3.300228
## [32,] 4.639808 -2.165905
## [33,] 1.289056 -4.468089
## [34,] 2.645605 -2.518130
## [35,] 2.333662 -2.673494
## [36,] 1.439703 -2.522436
## [37,] 2.677371 -4.374141
## [38,] 2.987012 -3.656608
## [39,] 1.159585 -1.568452
## [40,] 2.088574 -2.658771
## [41,] 3.866457 -2.475074
## [42,] 2.349272 -4.269171
## [43,] 3.448025 -2.312825
## [44,] 1.271517 -3.232274
## [45,] 4.432769 -3.584845
## [46,] 4.484413 -1.918888
## [47,] 3.049036 -3.211672
## [48,] 4.867017 -2.776078
## [49,] 2.817329 -2.631064
## [50,] 2.529022 -3.964889
## [51,] 2.106063 -5.056103
## [52,] 2.961445 -4.114520
## [53,] 1.375293 -2.475593
## [54,] 5.196439 -1.836739
## [55,] 3.751103 -4.130707
## [56,] 2.123831 -2.936833
## [57,] 3.298901 -3.883268
## [58,] 4.020068 -3.284051
## [59,] 2.976836 -2.521531
## [60,] 0.359280 -3.903877
```

```
plot(x)
```



Q. Can we use kmeans to cluster this data setting k 2 and nstart to 20

```
km <- kmeans(x, centers=2, nstart=20)
km
```

[illegible]

Q. How many points are in each cluster?

```
## [1] 30 30
```

Q. What ‘component’ of your result object details cluster assignment/membership?


```
## [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2 2
## [39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

What ‘component’ of your result object details cluster centers?

```
##           x           y
## 1 -3.147542  2.920132
## 2  2.920132 -3.147542
```

Q. Plot `x` colored by the `kmeans` cluster assignment and add cluster centers as blue points

```
plot(x, col= km$cluster)
points(km$centers, col="blue", pch=15, cex=2)
```




Analyze same data with `hclust()`

```
hc <- hclust (dist(x))
hc
```

```
##
## Call:
## hclust(d = dist(x))
##
## Cluster method   : complete
## Distance         : euclidean
## Number of objects: 60
```

There is a plot method for hclust result object

```
plot(hc)
```



To get our cluster membership vector we need to do a little more work. We have to “cut” the tree where we think it makes sense. For this we use `cutree()` function.

```
cutree(hc, h = 6)
```

```
## [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
## [39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```


You can also call `cutree()` setting `k = numner of grps/clusters you want`.

```
cutree(hc, k=2)
```

```
## [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2  
## [39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

Make our results plot

```
grps <- cutree(hc,k=2)  
plot(x, col=grps)
```



#PCA of UK food data

import data

```
url <- "https://tinyurl.com/UK-foods"  
x <- read.csv(url)
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
## [1] 17 5
```

## Checking your data

Preview first 6 rows

```
x
```

		X	England	Wales	Scotland	N.Ireland
##						
##	1	Cheese	105	103	103	66
##	2	Carcass_meat	245	227	242	267
##	3	Other_meat	685	803	750	586
##	4	Fish	147	160	122	93
##	5	Fats_and_oils	193	235	184	209
##	6	Sugars	156	175	147	139
##	7	Fresh_potatoes	720	874	566	1033
##	8	Fresh_Veg	253	265	171	143
##	9	Other_Veg	488	570	418	355
##	10	Processed_potatoes	198	203	220	187
##	11	Processed_Veg	360	365	337	334
##	12	Fresh_fruit	1102	1137	957	674
##	13	Cereals	1472	1582	1462	1494
##	14	Beverages	57	73	53	47
##	15	Soft_drinks	1374	1256	1572	1506
##	16	Alcoholic_drinks	375	475	458	135
##	17	Confectionery	54	64	62	41

Row-names are incorrectly set as the first column of our x data fram. Use minus indexing as one option.

```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

		England	Wales	Scotland	N.Ireland
##					
##	Cheese	105	103	103	66
##	Carcass_meat	245	227	242	267
##	Other_meat	685	803	750	586
##	Fish	147	160	122	93
##	Fats_and_oils	193	235	184	209
##	Sugars	156	175	147	139

There is a problem here because if we keep inputing the function it takes a column away each time.

Another method:

```
x <- read.csv(url, row.names=1)
head(x)
```

		England	Wales	Scotland	N.Ireland
##					
##	Cheese	105	103	103	66
##	Carcass_meat	245	227	242	267
##	Other_meat	685	803	750	586
##	Fish	147	160	122	93
##	Fats_and_oils	193	235	184	209
##	Sugars	156	175	147	139

```
dim(x)
```


```
## [1] 17 4
```

Now we have correct dimensions > Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

The second option is the better option because it doesn’t keep taking away columns.

Plot the data

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```




Q3: Changing what optional argument in the above `barplot()` function results in the following plot?

Set beside to false


```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```





Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(10), pch=16)
```



Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

Can't really tell the difference at this point the way the data is presented.


Use PCA Funtion

```
# Use the prcomp() PCA function
pca <- prcomp(t(x))
summary(pca)
```

```
## Importance of components:
##              PC1      PC2      PC3      PC4
## Standard deviation 324.1502 212.7478 73.87622 4.189e-14
## Proportion of Variance 0.6744 0.2905 0.03503 0.000e+00
## Cumulative Proportion 0.6744 0.9650 1.00000 1.000e+00
```

Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
#Plot PC1 vs PC2
colors <- c("red", "Orange", "Green", "Blue")
plot(pca$x[,1], pca$x[,2])
text(pca$x[,1], pca$x[,2], colnames(x), col= colors)
```



Variable Loadings!

```
# Focus on PC1 as it accounts for > 90% of variance  
par(mar=c(10, 3, 0.35, 0))  
barplot(pca$rotation[,1], las=2)
```

